

IoT Big Data Processing

Distributed Streaming Classification

Albert Bifet(@abifet)




Motivation

- Datasets already stored on clusters
 - Don't want to move everything to single powerful machine
- Clusters ubiquitous and cheap (e.g., see TOP500), supercomputers expensive and monolithic
 - Clusters easily shared, leverage economy of scale
- Largest problem solvable by single machine constrained by hardware
 - How fast can you read from disk or network

Distributed Stream Processing Engines

A Tale of two Tribes

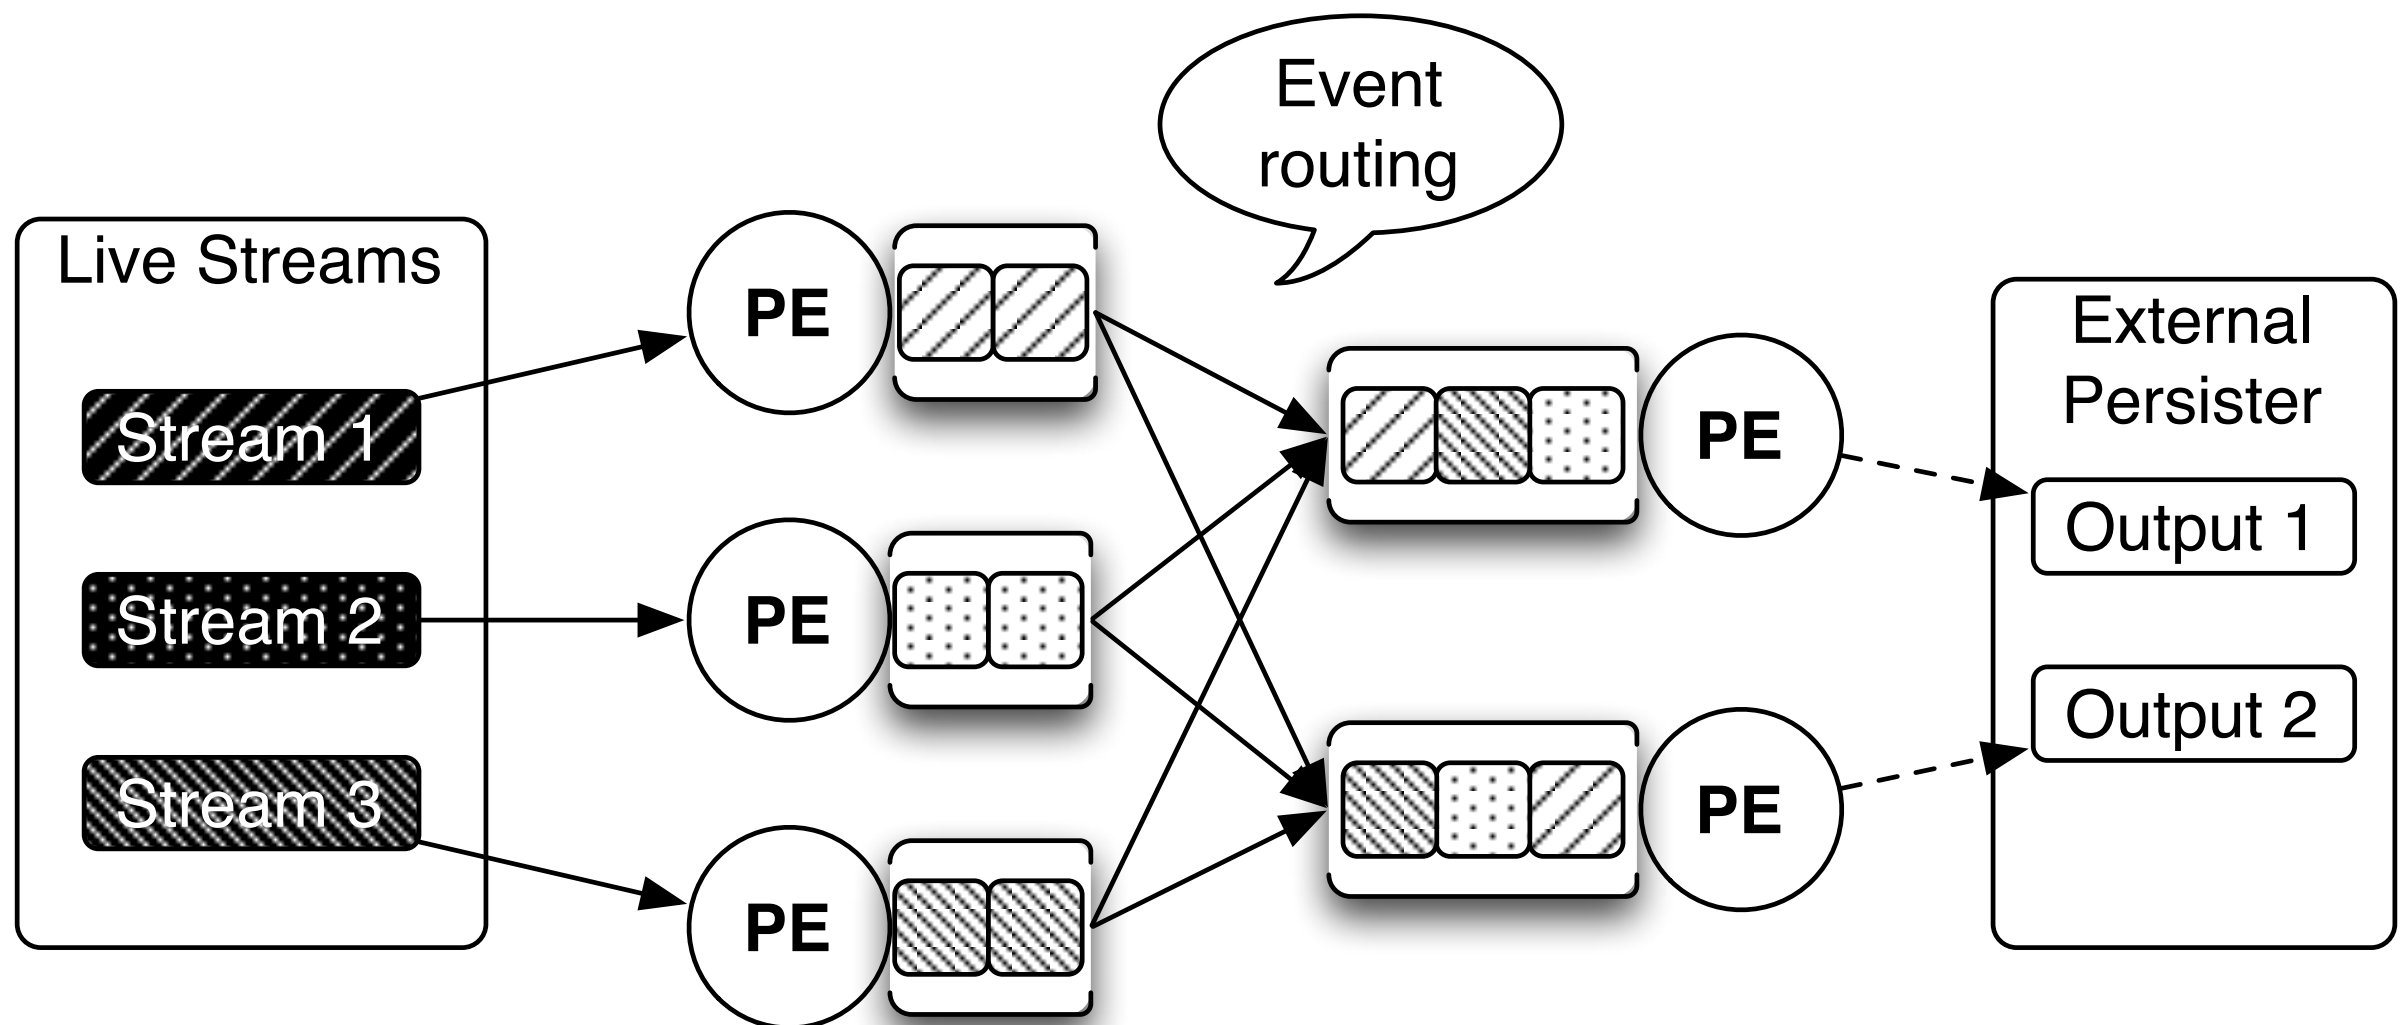
M. Stonebraker U. Çetintemel: “One Size Fits All’: An Idea Whose Time Has Come and Gone”. ICDE ’05



SPE Evolution

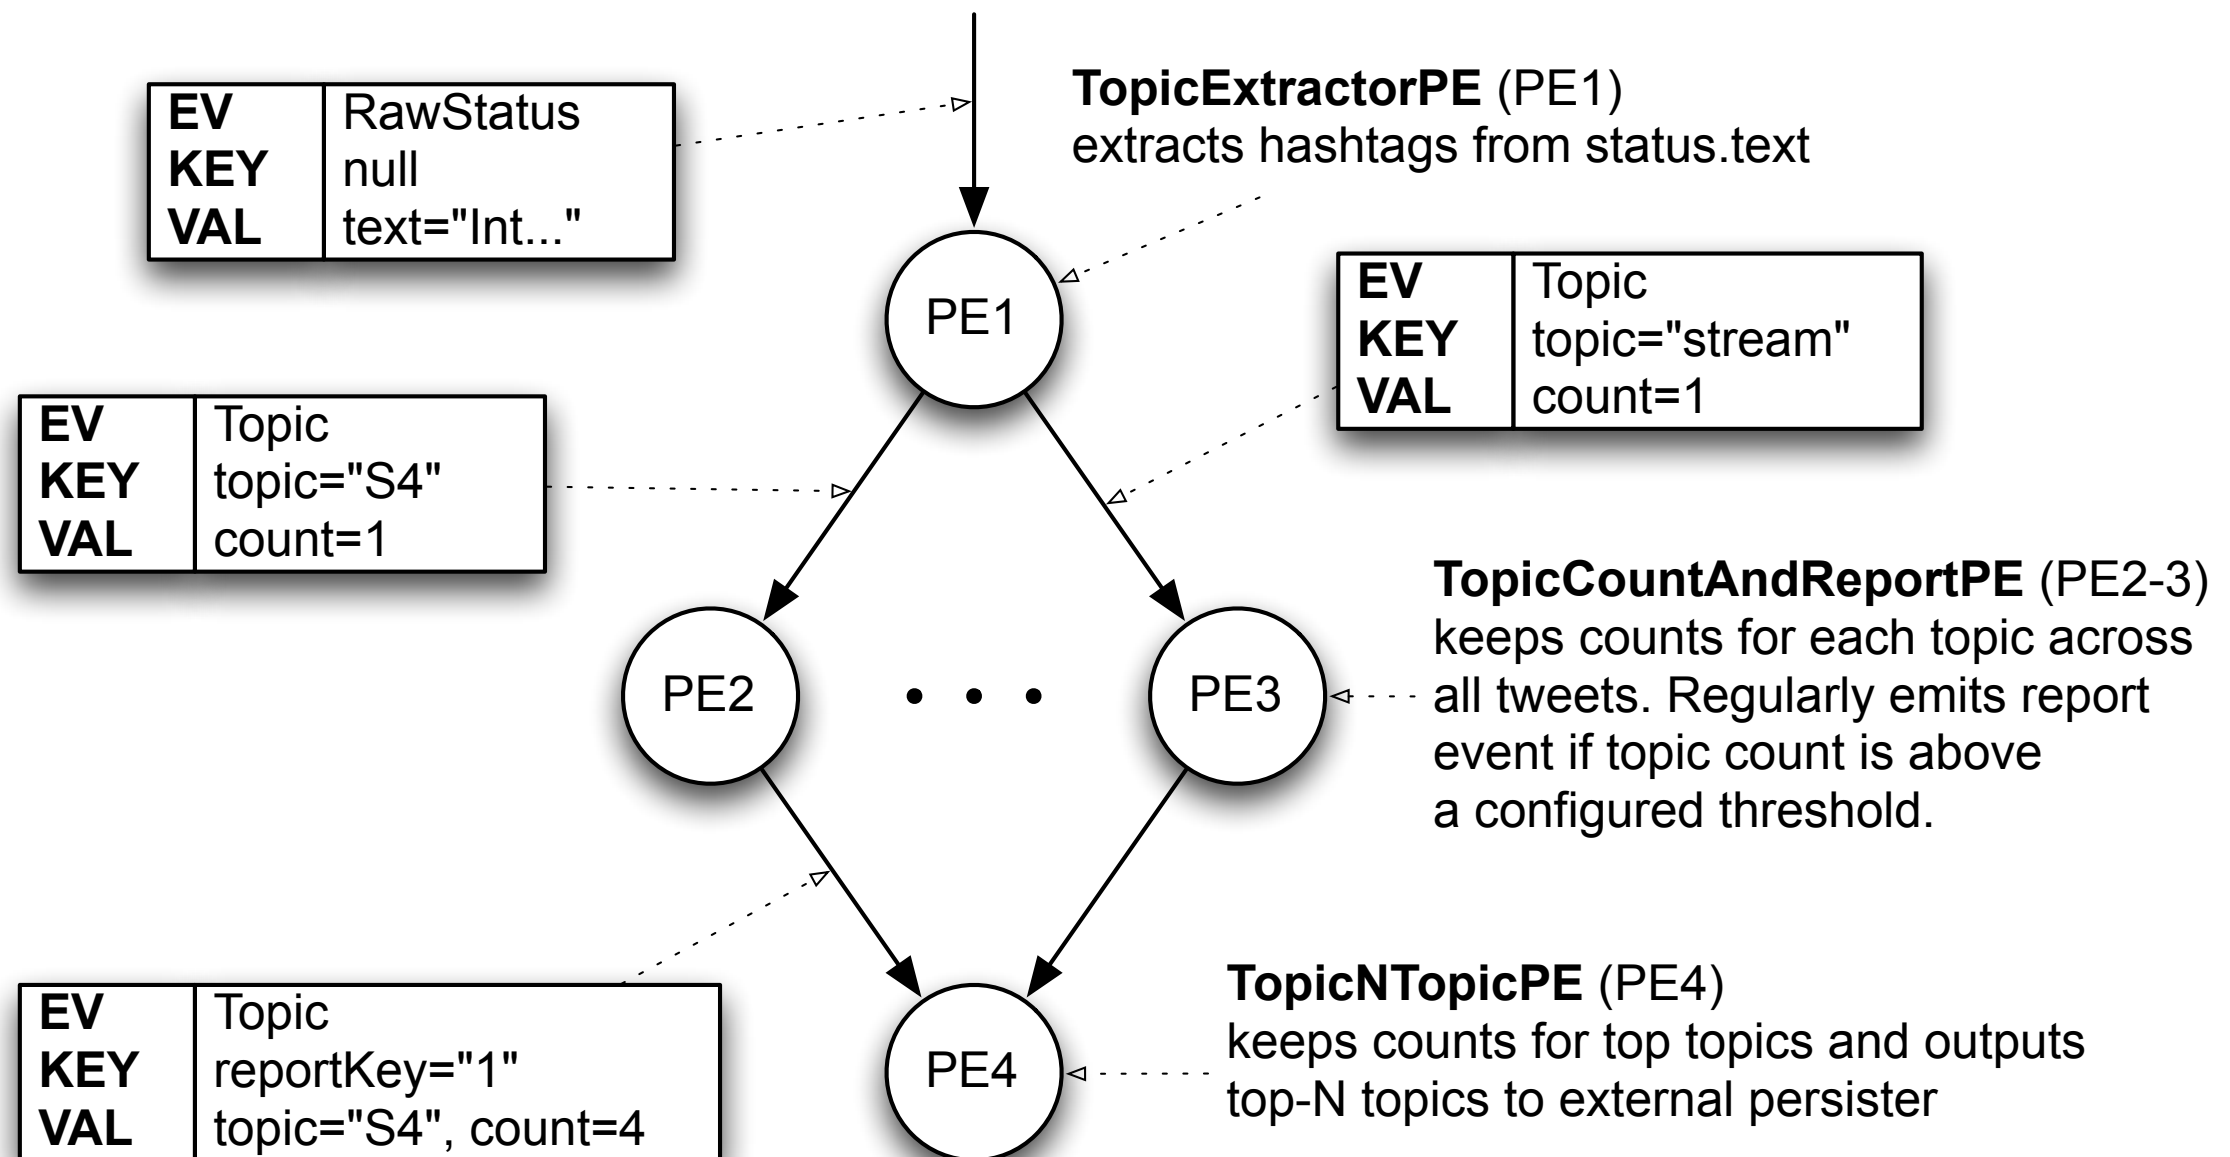
1 st generation	–2003	Aurora	Abadi et al., “Aurora: a new model and architecture for data stream management,” VLDB Journal, 2003
	–2004	STREAM	Arasu et al., “STREAM: The Stanford Data Stream Management System,” Stanford InfoLab, 2004.
2 nd generation	–2005	Borealis	Abadi et al., “The Design of the Borealis Stream Processing Engine,” in CIDR ’05
	–2006	SPC	Amini et al., “SPC: A Distributed, Scalable Platform for Data Mining,” in DMSSP ’06
	–2008	SPADE	Gedik et al., “SPADE: The System S Declarative Stream Processing Engine,” in SIGMOD ’08
3 rd generation	–2010	S4	Neumeyer et al., “S4: Distributed Stream Computing Platform,” in ICDMW ’10
	–2011	Storm	http://storm.apache.org
	–2013	Samza	http://samza.incubator.apache.org

Actors Model




S4 Example

status.text: "Introducing #S4: a distributed #stream processing system"




Groupings

- Key Grouping
(hashing)
- Shuffle Grouping
(round-robin)
- All Grouping
(broadcast)




Groupings

- **Key Grouping (hashing)**
- Shuffle Grouping (round-robin)
- All Grouping (broadcast)




Groupings

- Key Grouping
(hashing)
- **Shuffle Grouping
(round-robin)**
- All Grouping
(broadcast)



Groupings

- Key Grouping (hashing)
- Shuffle Grouping (round-robin)
- **All Grouping (broadcast)**




Classification

Hadoop AllReduce

A. Agarwal, O. Chapelle, M. Dudík, J. Langford: “A Reliable Effective Terascale Linear Learning System”. JMLR (2014)

- MPI AllReduce on MapReduce
- Parallel SGD + L-BFGS
- Aggregate + Redistribute
 - Each node computes partial gradient
 - Aggregate (sum) complete gradient
 - Each node gets updated model
- Hadoop for data locality (map-only job)





AllReduce



Reduction Tree

Upward = Reduce

Downward = Broadcast (All)

Parallel Decision Trees


- Which kind of parallelism?
 - Task
 - Data
 - Horizontal
 - Vertical



Decision Tree

- Each node tests a features
- Each branch represents a value
- Each leaf assigns a class
- Greedy recursive induction
 - Sort all examples through tree
 - x_i = most discriminative attribute
 - New node for x_i , new branch for each value, leaf assigns majority class
 - Stop if no error | limit on #instances

Car deal?



Very Fast Decision Tree

Pedro Domingos, Geoff Hulten: "Mining high-speed data streams". KDD '00

- AKA, Hoeffding Tree
- A small sample can often be enough to choose a near optimal decision
- Collect sufficient statistics from a small set of examples
- Estimate the merit of each alternative attribute
- Choose the sample size that allows to differentiate between the alternatives

Leaf Expansion

- When should we expand a leaf?
- Let x_1 be the most informative attribute, x_2 the second most informative one
- Is x_1 a stable option?
- Hoeffding bound
 - Split if $G(x_1) - G(x_2) > \varepsilon = \sqrt{\frac{R^2 \ln(1/\delta)}{2n}}$

HT Induction

HT(*Stream*, δ)

- 1 ▷ Let HT be a tree with a single leaf(root)
- 2 ▷ Init counts n_{ijk} at root
- 3 **for** each example (x, y) in Stream
- 4 **do** HTGROW((x, y) , HT, δ)

HTGROW((x, y) , HT, δ)

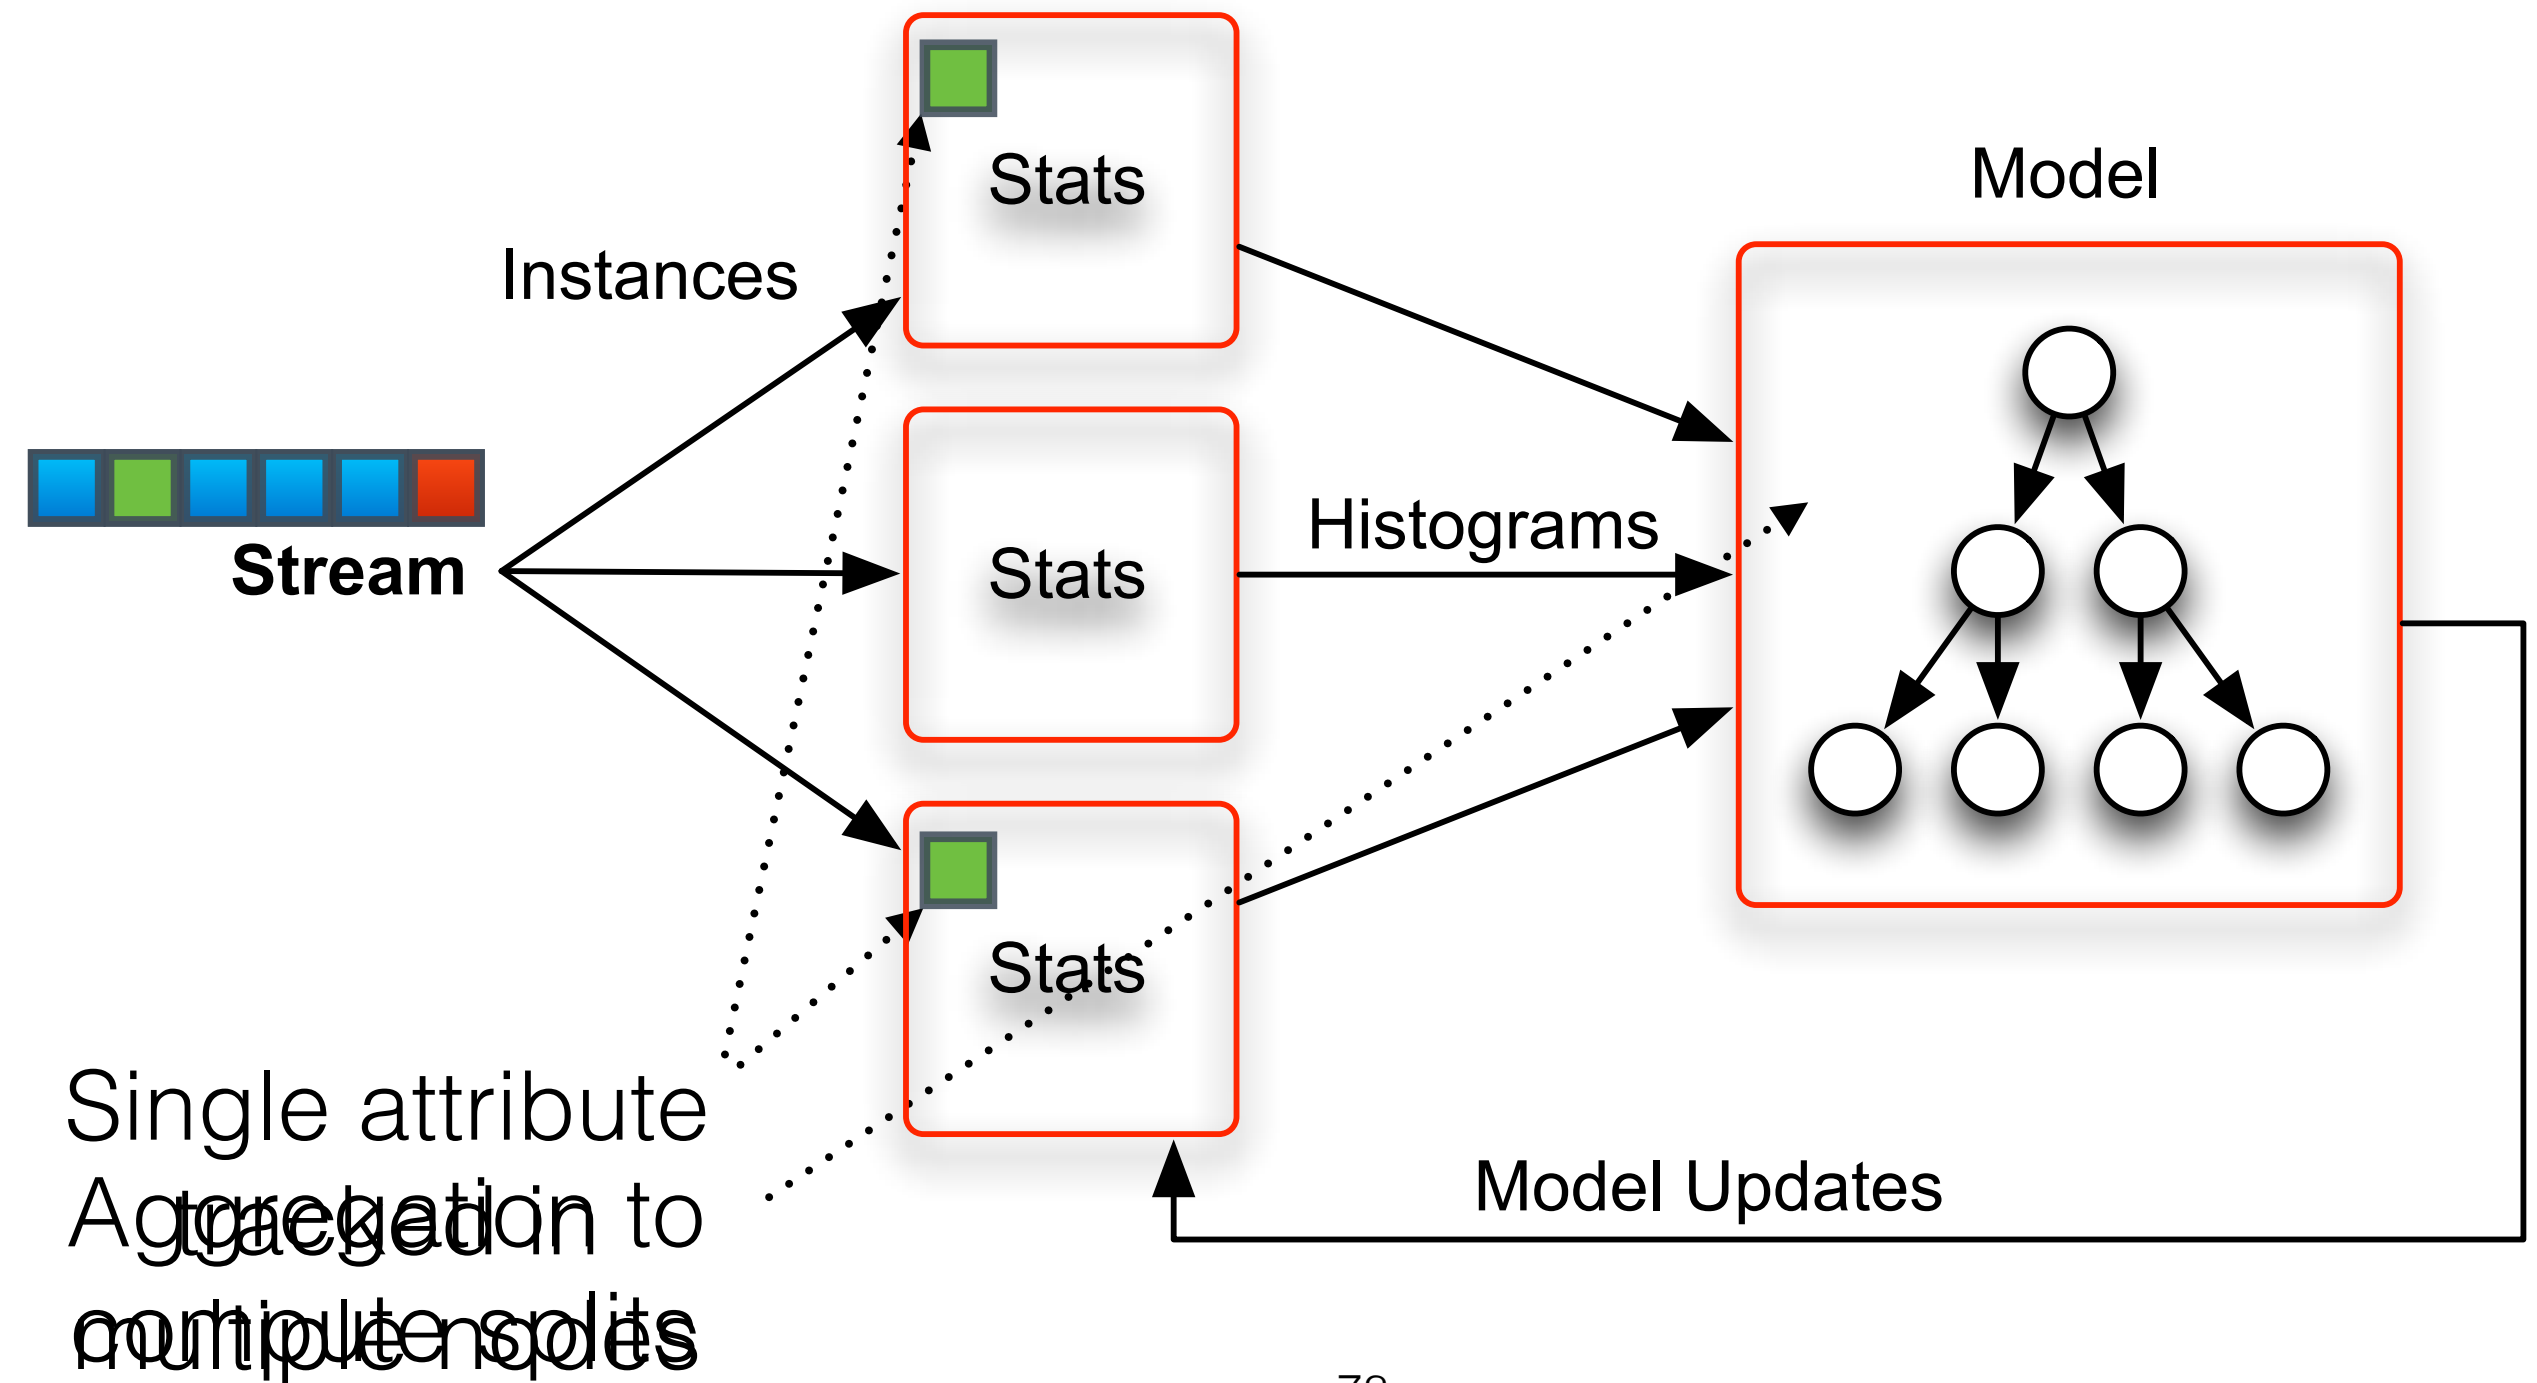
- 1 ▷ Sort (x, y) to leaf l using HT
- 2 ▷ Update counts n_{ijk} at leaf l
- 3 **if** examples seen so far at l are not all of the same class
- 4 **then** ▷ Compute G for each attribute
- 5 **if** $G(\text{Best Attr.}) - G(\text{2nd best}) > \sqrt{\frac{R^2 \ln 1/\delta}{2n}}$
- 6 **then** ▷ Split leaf on best attribute
- 7 **for** each branch
- 8 **do** ▷ Start new leaf and initialize counts

Properties

- Number of examples to expand node depends only on Hoeffding bound (ϵ decreases with \sqrt{n})
- Low variance model (stable decisions with statistical support)
- Low overfitting (examples processed only once, no need for pruning)
- Theoretical guarantees on error rate with high probability
 - Hoeffding algorithms asymptotically close to batch learner.
Expected disagreement δ/p (p = probability instance falls into a leaf)
- Ties: broken when $\epsilon < \tau$ even if $\Delta G < \epsilon$


Horizontal Partitioning

Y. Ben-Haim, E. Tom-Tov: "A Streaming Parallel Decision Tree Algorithm". JMLR (2010)



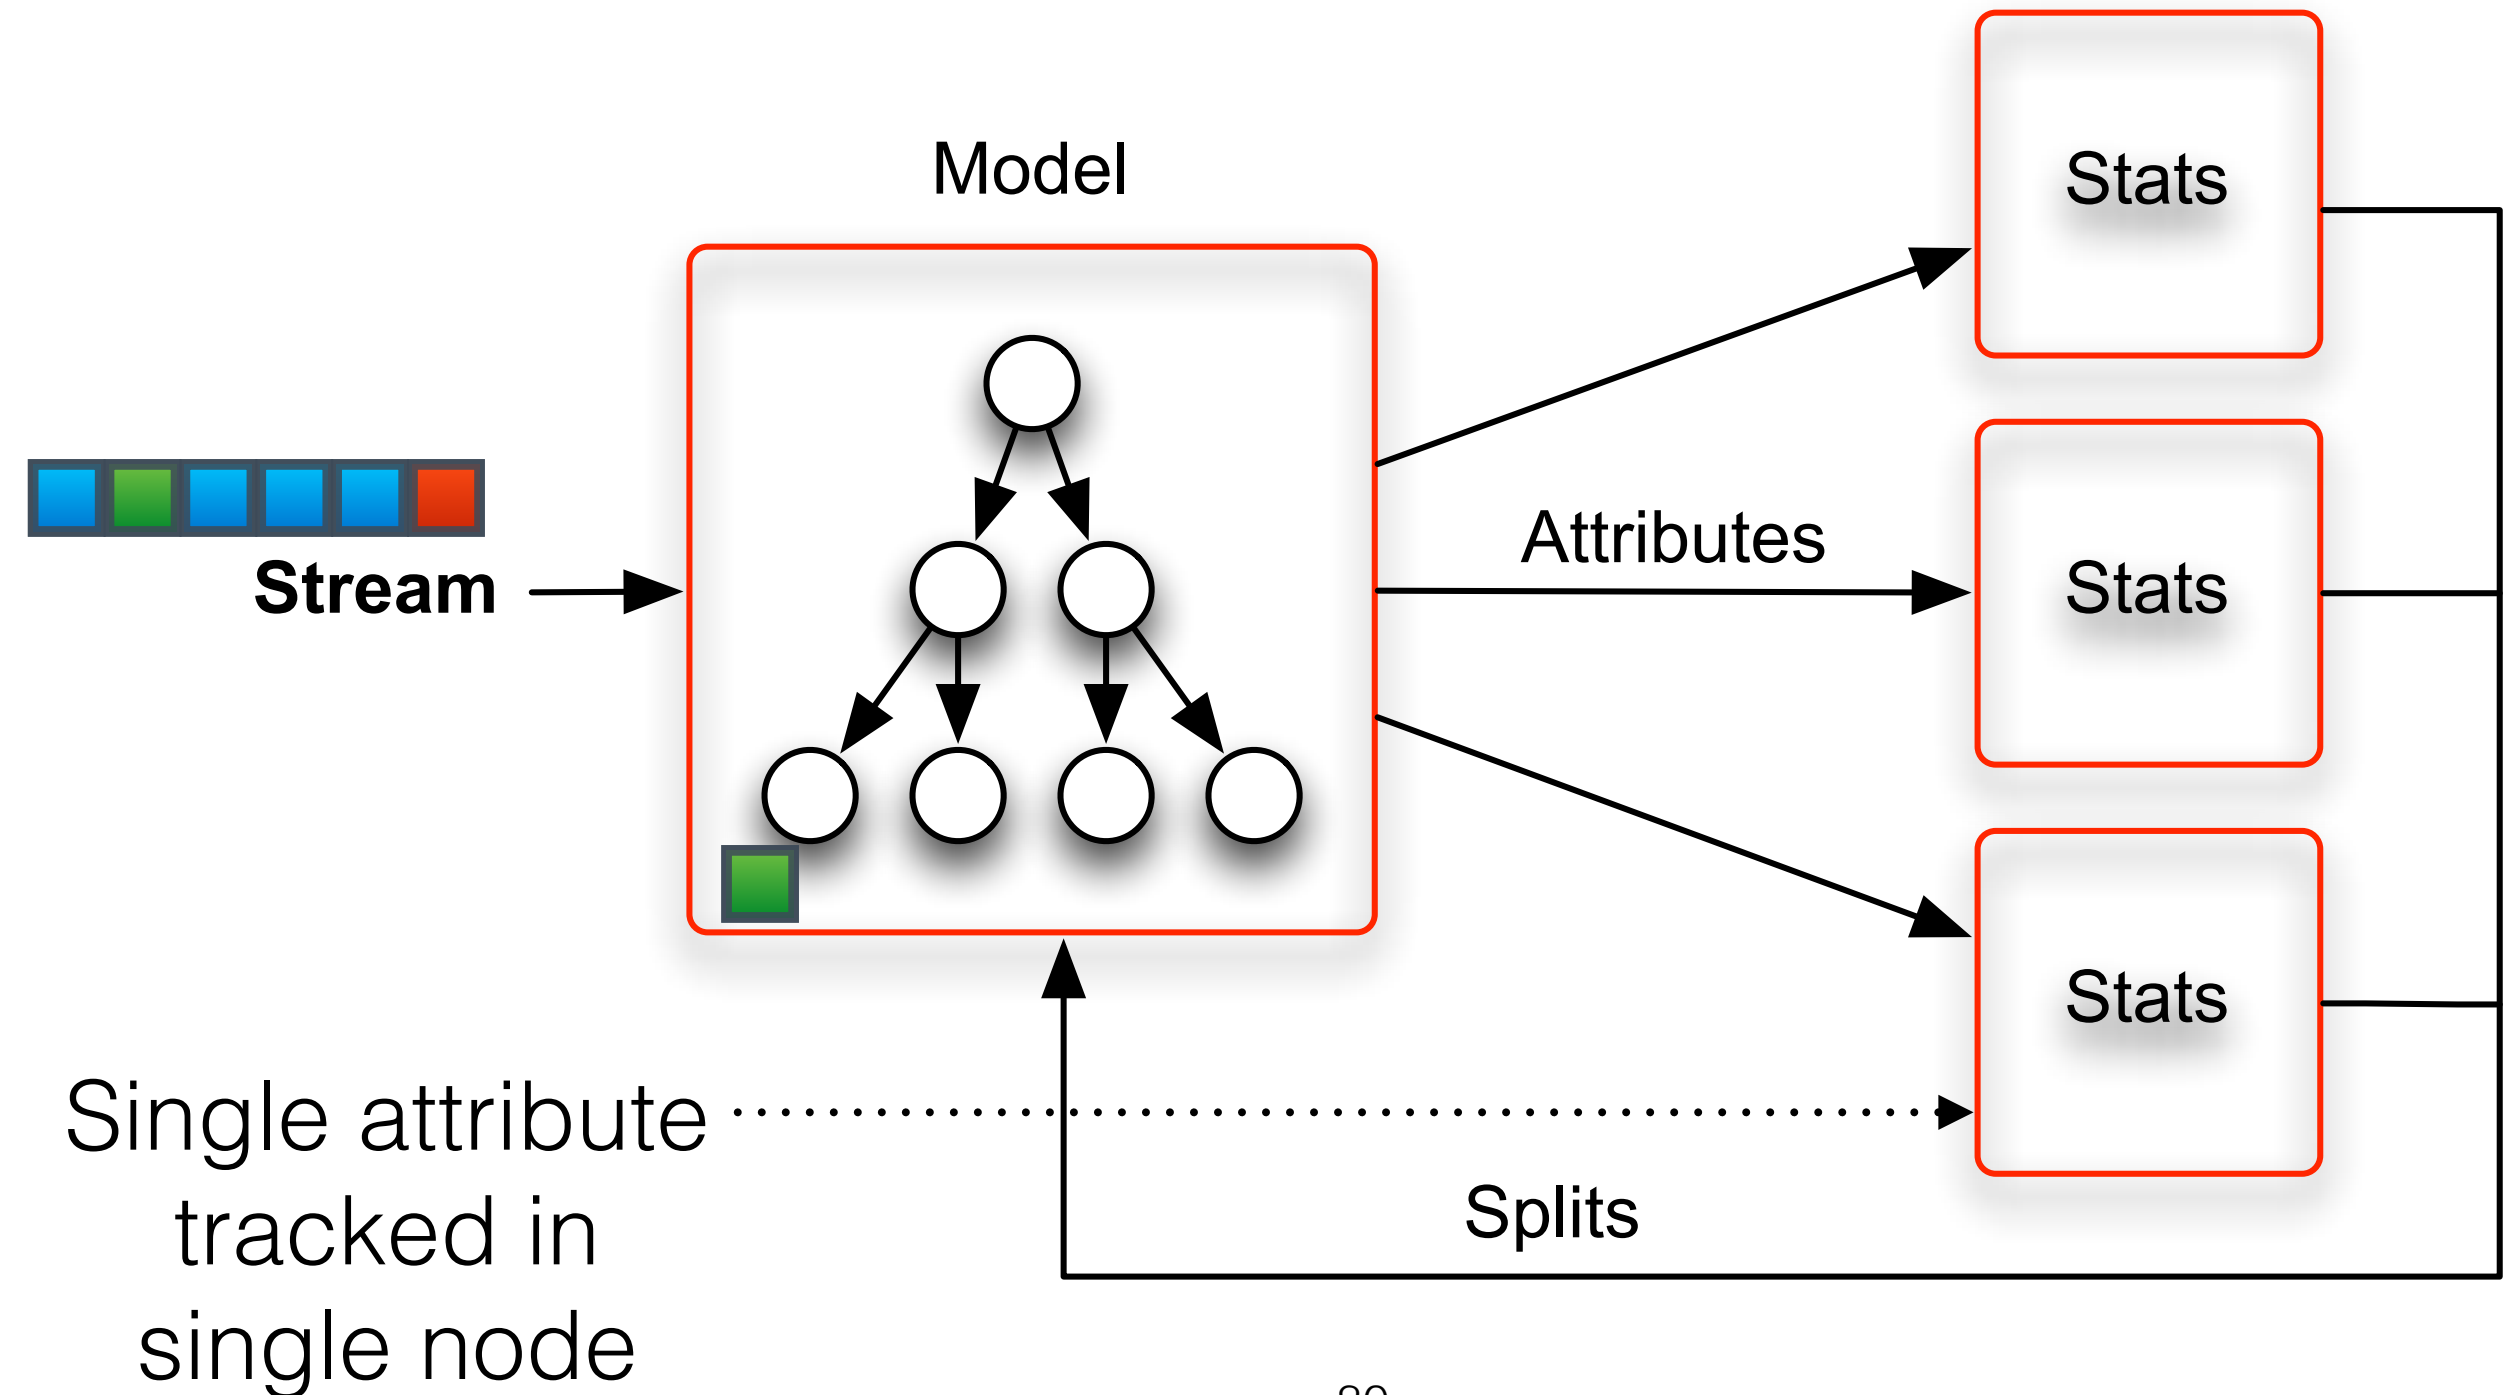
Hoeffding Tree Profiling

Training time for
100 nominal +
100 numeric
attributes

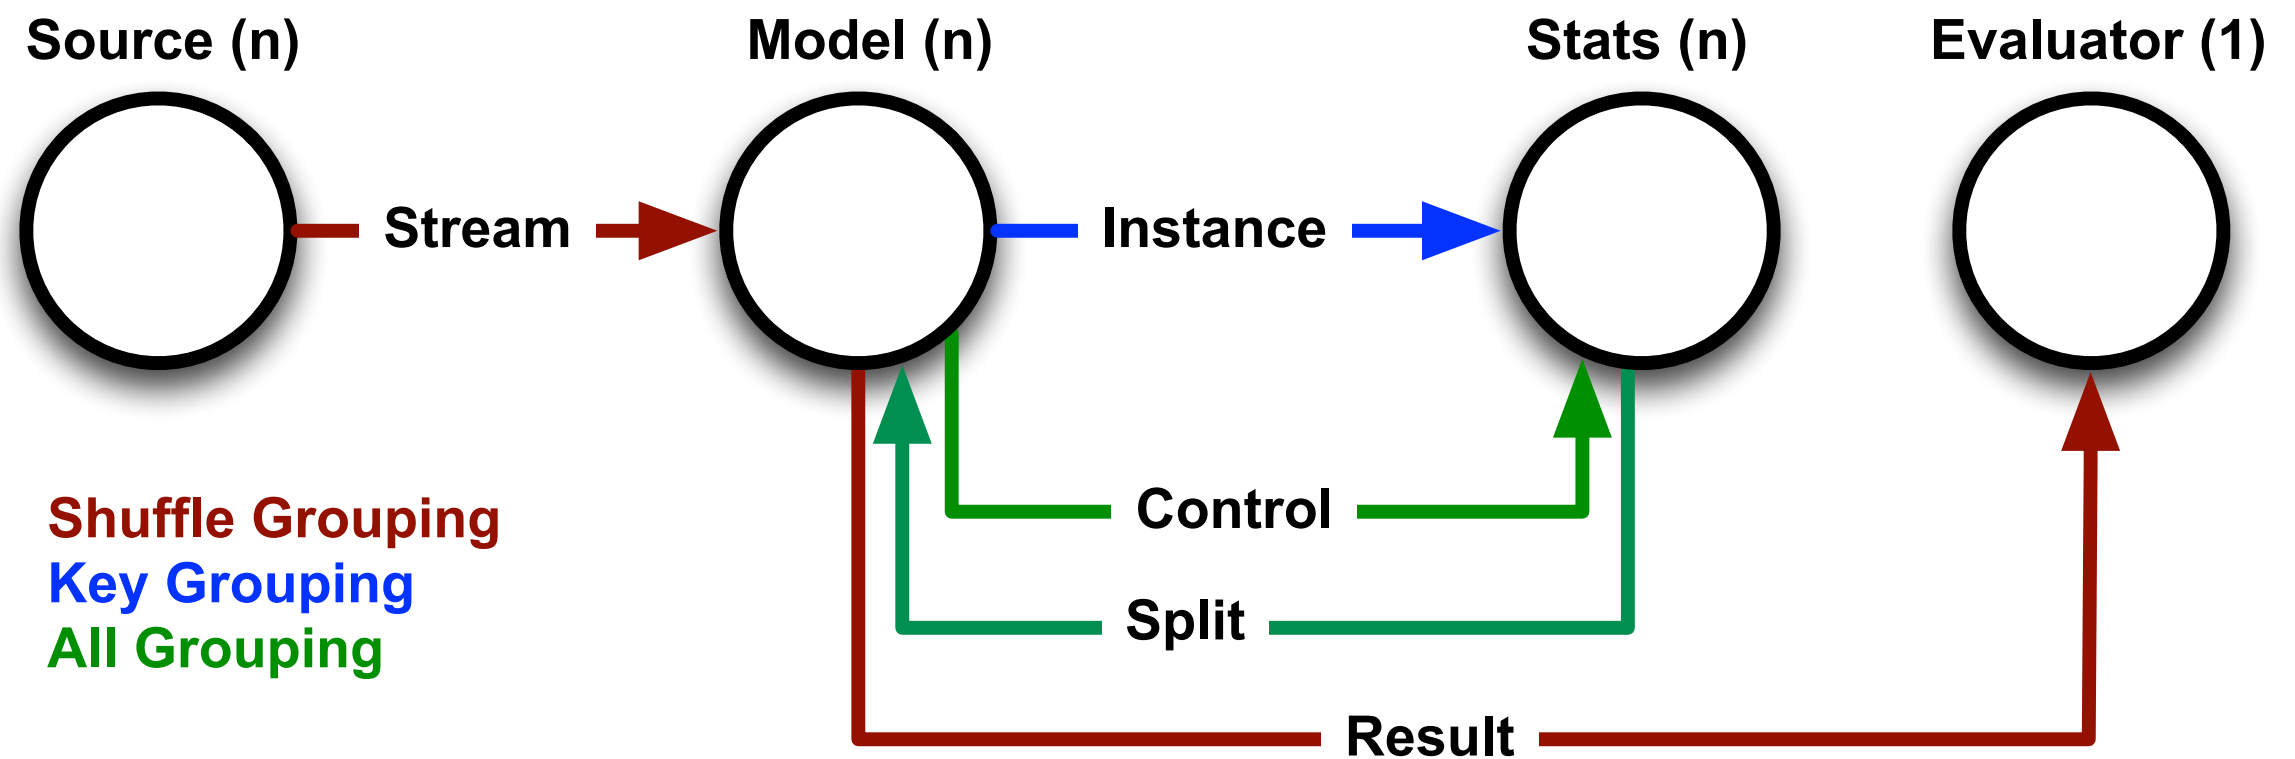


Vertical Partitioning

A. Murdopo, A. Bifet, G. De Francisci Morales, N. Kourtellis: "VHT: Vertical Hoeffding Tree". Working paper (2014)



Vertical Hoeffding Tree

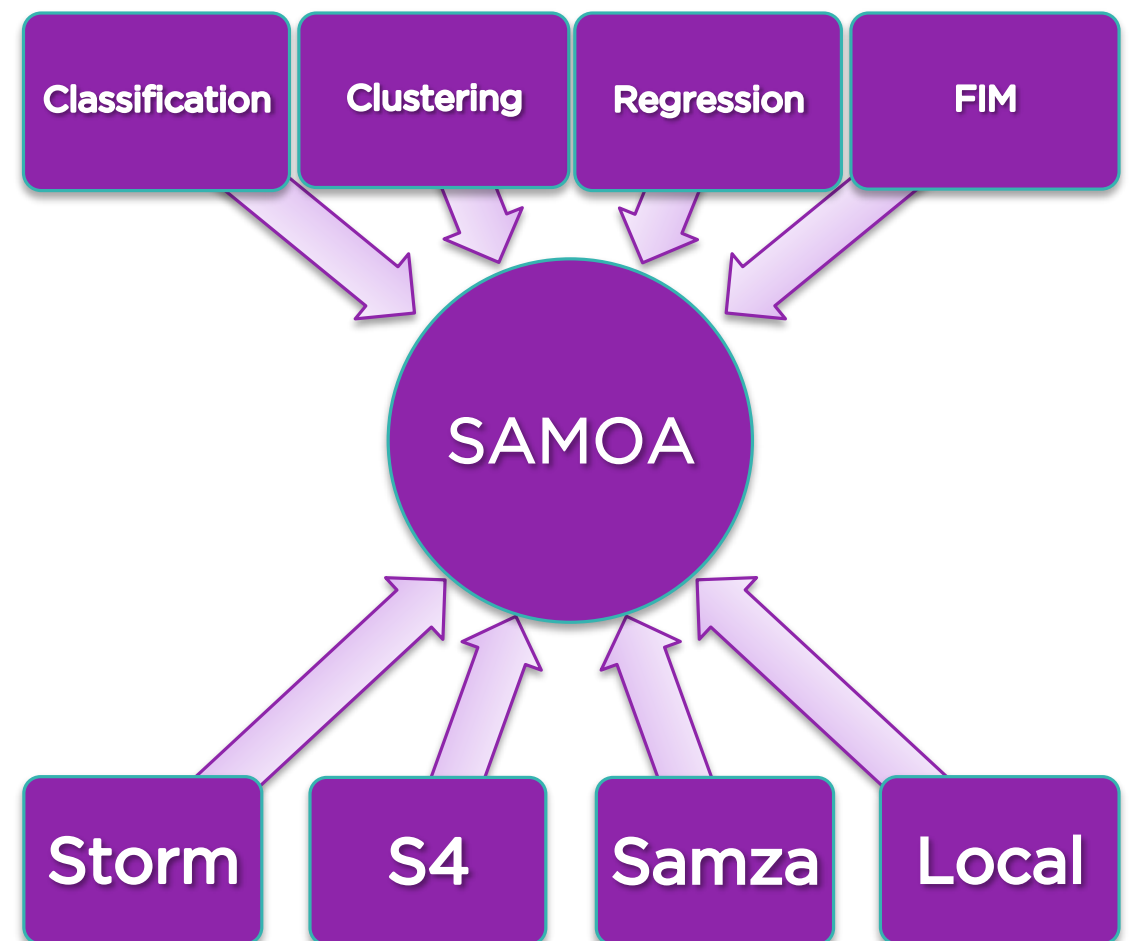
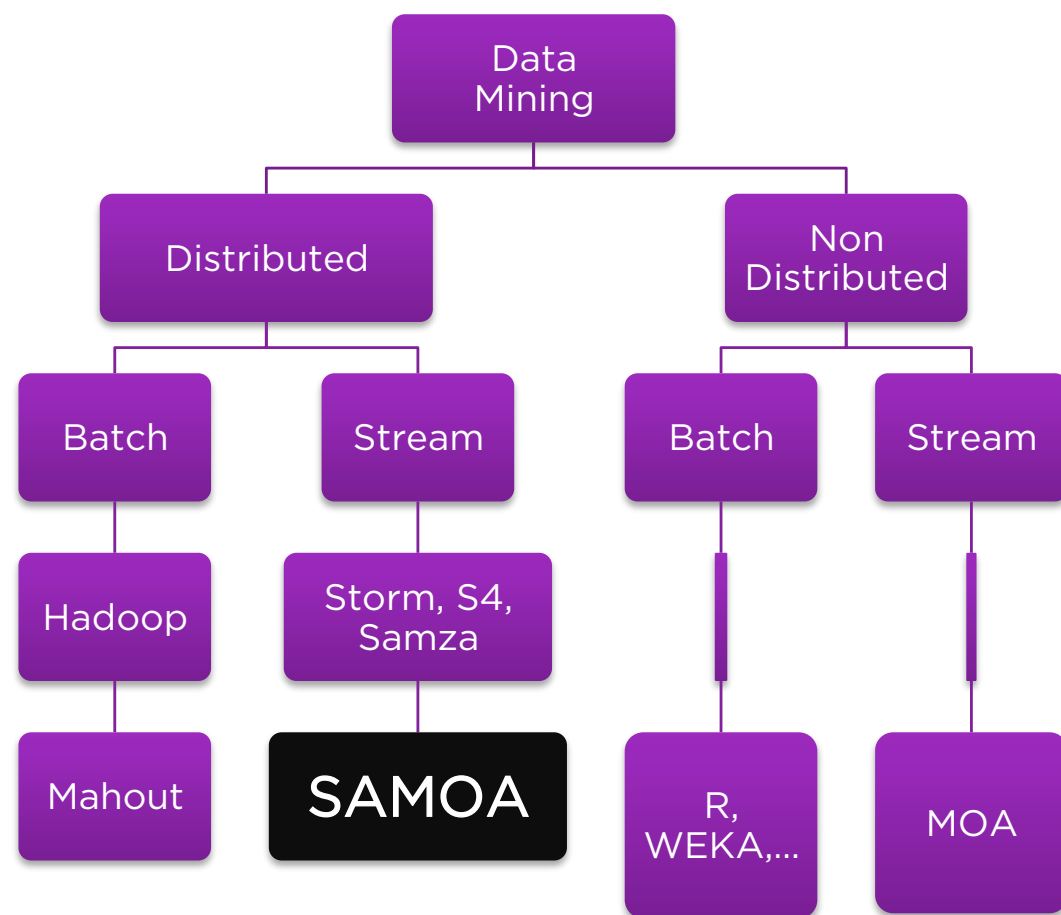


Advantages of Vertical Parallelism

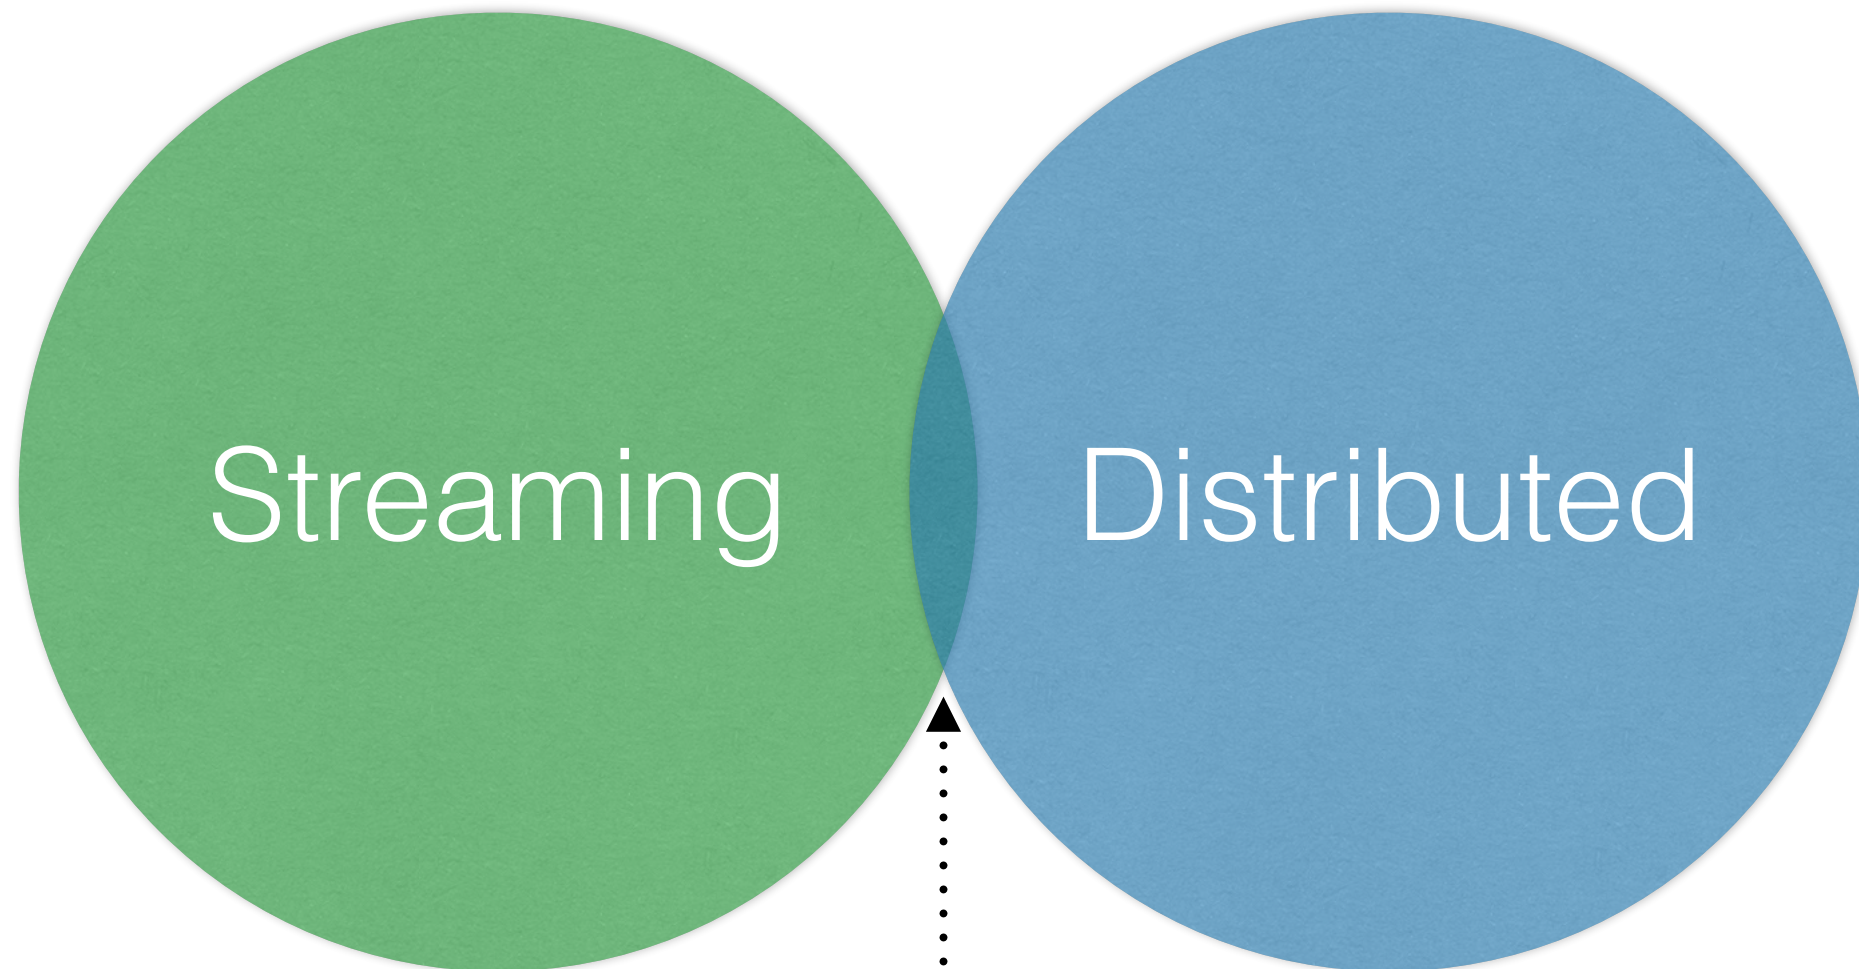
- High number of attributes => high level of parallelism (e.g., documents)
- vs. task parallelism
 - Parallelism observed immediately
- vs. horizontal parallelism
 - Reduced memory usage (no model replication)
 - Parallelized split computation

SAMOA

G. De Francisci Morales, A. Bifet: "SAMOA: Scalable Advanced Massive Online Analysis". JMLR (2014)



Vision



Big Data Stream Mining